



SNOWFLAKE BUSINESS ANALYST WORKSHOP

State of Colorado

October 2025

Agenda

1. Snowflake Overview
2. Hands-on Lab: Data Exploration, Querying, & Analytics
3. Q&A and Wrap-up



Snowflake Overview

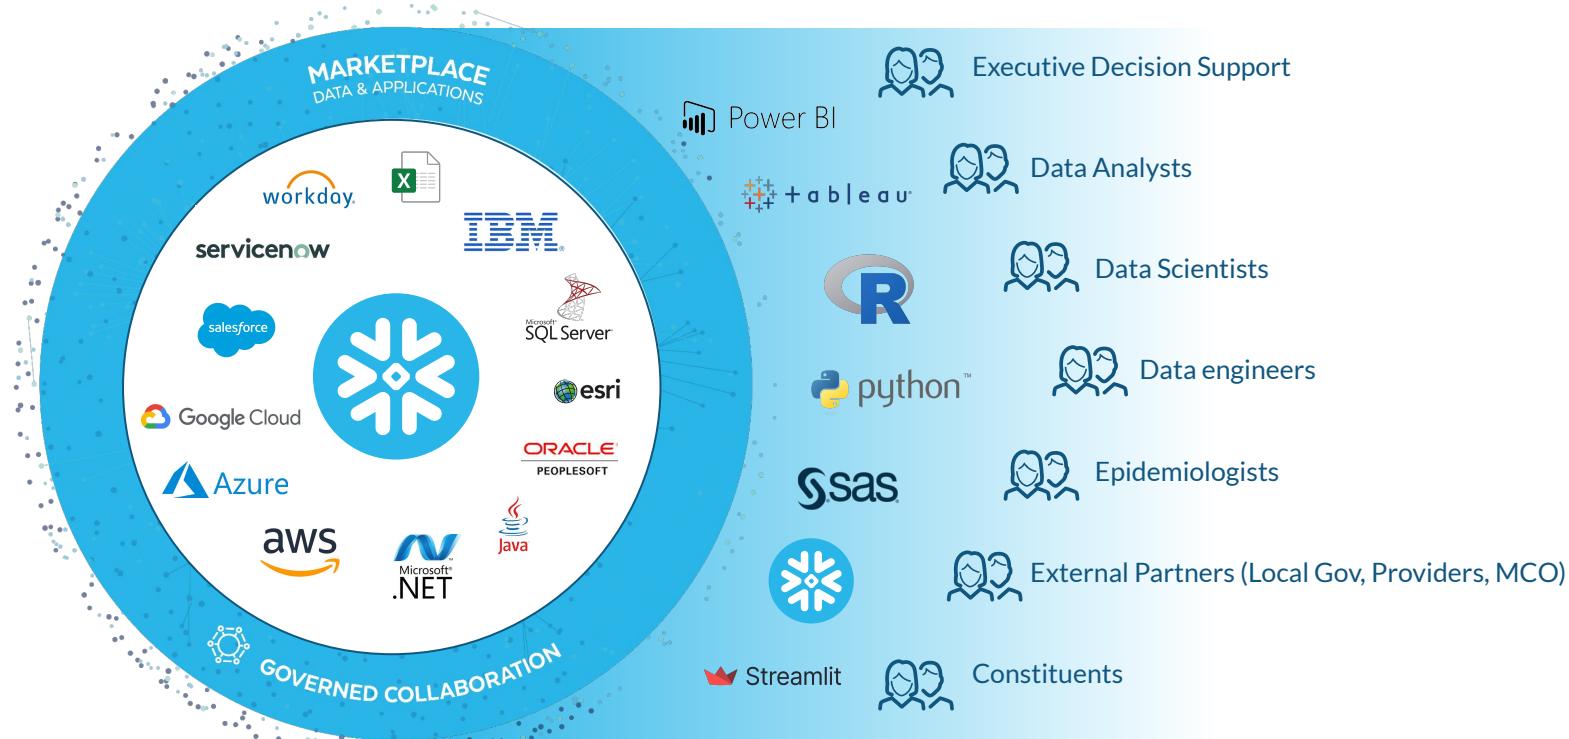
1:00 PM - 1:30 PM



© 2025 Snowflake Inc. All Rights Reserved

A Single Source of Truth

Shatter data silos and democratize insights to every team



One Platform. Any Architectural Pattern.

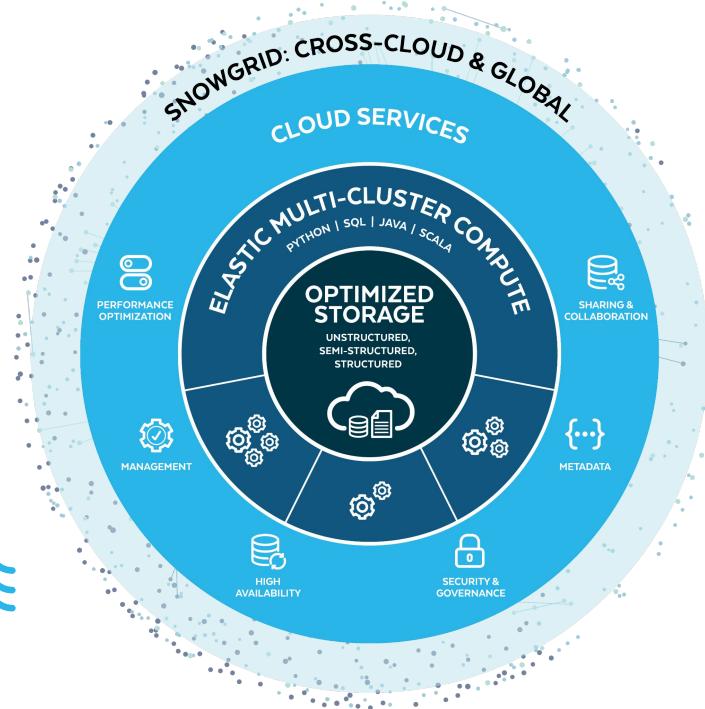
Data Warehouse

Centralized analytics for structured business data



Data Lakehouse

Transactional data lake for unified analytics, AI/ML, collaborative workloads



Data Lake

Unlimited storage for versatile data types and workloads

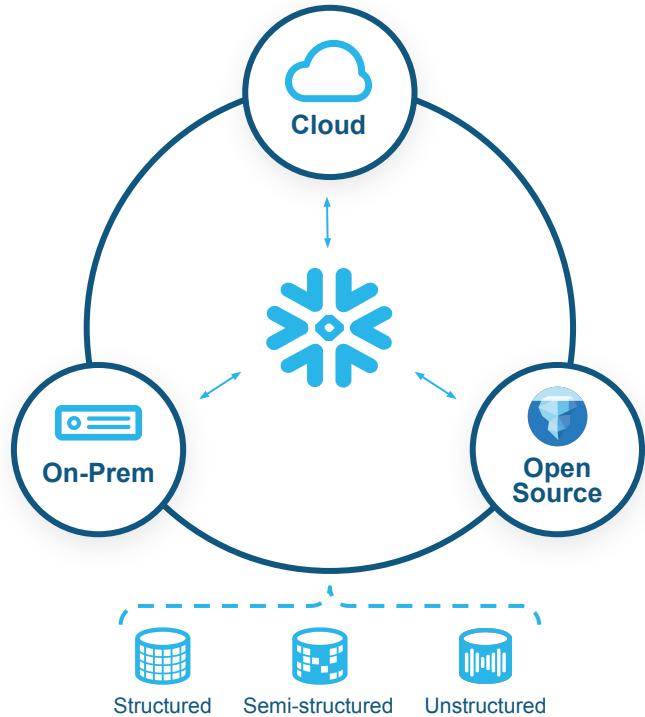


Data Mesh / Fabric

Distributed and governed, domain-oriented collaboration



Optimized Storage



Un-siloed Access to Your Data

Unstructured, semi-structured, and structured data together with near-infinite scale.

Easily Manage Data at Scale

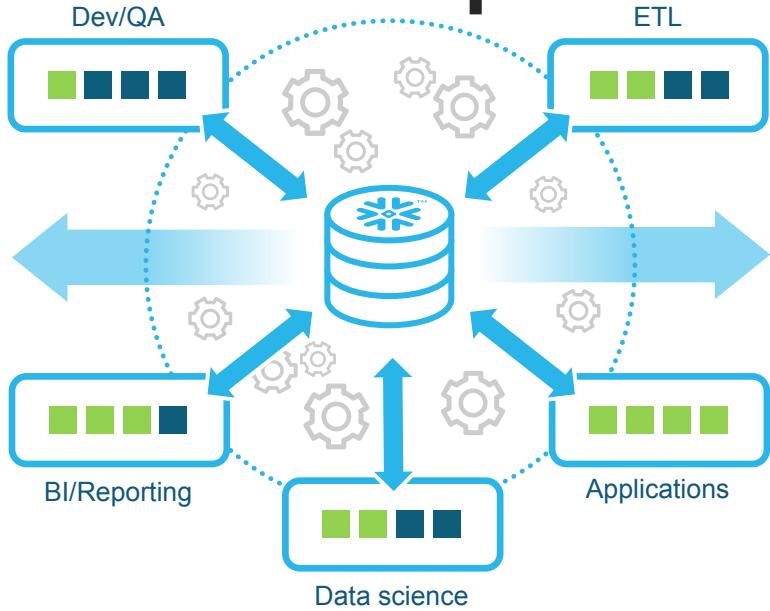
Fast and efficient access, optimized compression, and secure data - all automated.

Flexibility & Interoperability

Work with data on-premises or in open table formats* to remove lock-in and adapt to any architectural pattern.



Elastic Multi-Cluster Compute



One Engine for Every Workload

Power streaming pipelines, analytics, AI, interactive apps, and more through a single engine with flexible compute options including GPUs.

Leading Performance & Concurrency

Fast, reliable performance for virtually all users and jobs with no tuning or contention, delivered through isolated compute.

Familiar Languages

Work in SQL, Python, or Java. Run your preferred libraries with Snowpark. All without moving data.

Flexible Development

Use native development interfaces, bring your own IDE, or leverage popular third-party tools.





Cloud Services



Snowflake Managed

Governance Controls

AI Models & LLMs

Maintenance & Tuning

Performance Enhancements

Networking & Encryption

Administration & Availability

Fully Managed Service

Automate costly and complex operations to reduce overhead and improve efficiency. Transparent performance enhancements automatically applied via releases.

Unified, Built-In Governance

Snowflake Horizon delivers unified compliance, security, privacy, interoperability, and access capabilities, without additional configurations or protocols.

Accessible, Built-In AI

Snowflake Cortex* enables secure access to industry-leading AI models, LLMs, and vector search functionality, with no infrastructure to manage.



Near-Zero Management - It Just Works

	Task	Ordinary on-prem DW/DL	Ordinary cloud DW/DL	Snowflake
Infrastructure 	Hardware	You	Cloud vendor	Built-in
	Software	Distribution	Cloud vendor	Built-in
	Hardware cluster setup	You	You	Built-in
	Software provisioning	Tools	Tools	Built-in
Data & service protection 	Data protection & retention	Platform or Add-on	Platform or Add-on	Cloud Data Platform as a Service
	Node failure protection & recovery	You	You	
	Disaster recovery	You	You	
	Service monitoring & alerting	You	You	
Security 	Physical security	You	Cloud vendor	Cloud Data Platform as a Service
	Deployment security	You	You	
	Security monitoring	You	You	
Database management & tuning 	Compute scaling	You	You	Cloud Data Platform as a Service
	Index management	You	You	
	Data partitioning	You	You	
	Metadata & statistics management	You	You	
	Query optimization	You	You	

-  NO Infrastructure
-  NO Tuning
-  NO Optimization
-  NO Indexing
-  NO Storage Worries
-  NO Vacuuming
-  NO Partitioning
-  NO Required Sorting
-  NO Workload Mgmt.
-  NO Manual Backups

Protecting Your Data in Snowflake

End-to-End Encryption

Always-encrypted client communications, plus integration with cloud provider private networking



Fully Encrypted Storage

Data at rest is always encrypted while handled by the Snowflake drivers and systems



Strong Authentication

Built in multi-factor, integration with your federated SSO, easy user management



Full Auditing

Track every login, every transaction, every data transfer, and export to your security tools



Role-Based Access Control

All objects, actions, and even compute usage can be controlled with roles



Recovery

We give you options to ensure your data can be recovered in case of an accident or worse

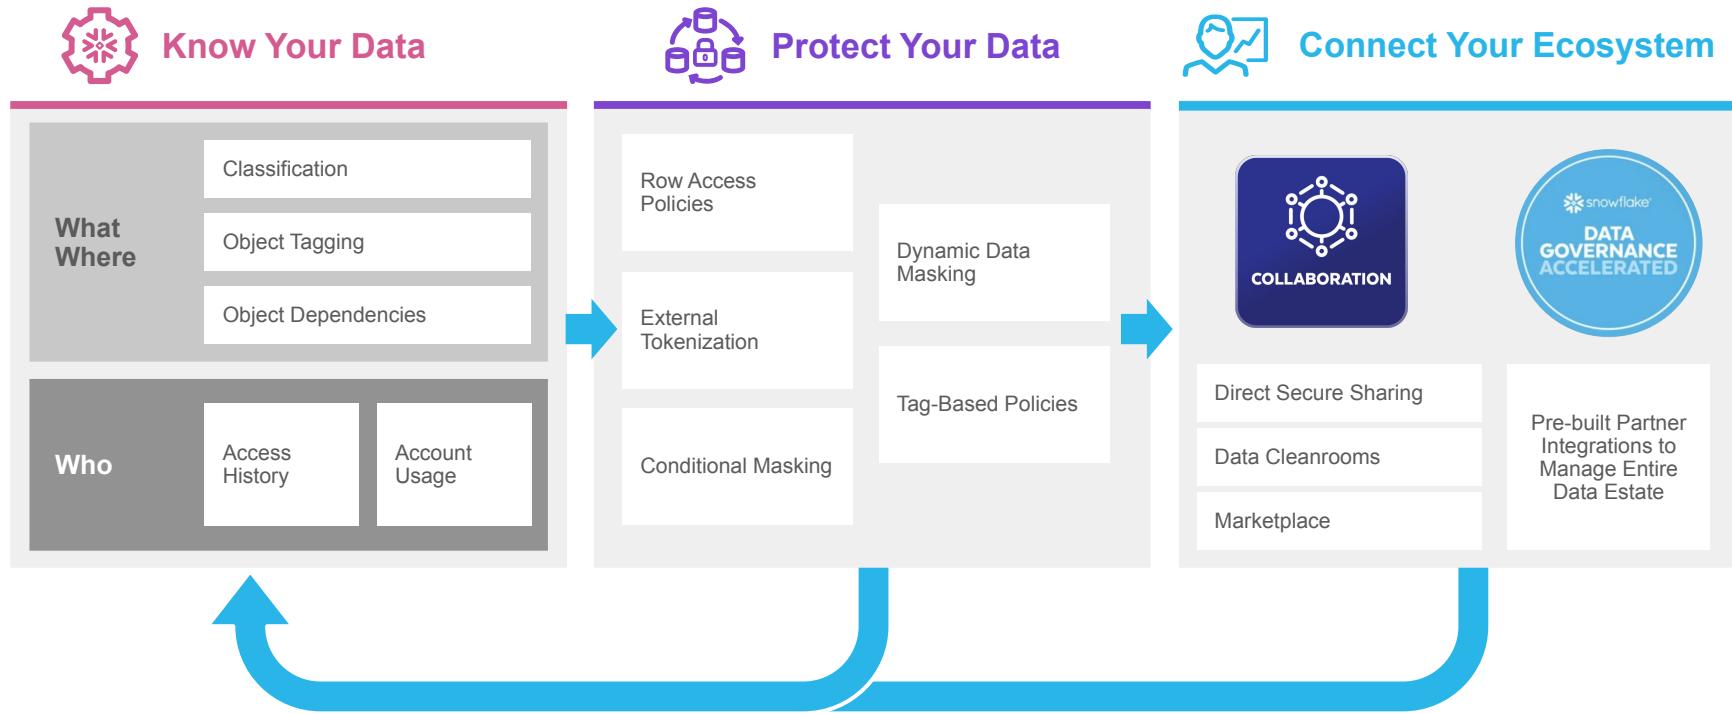


[Snowflake Security Product Documentation](#)

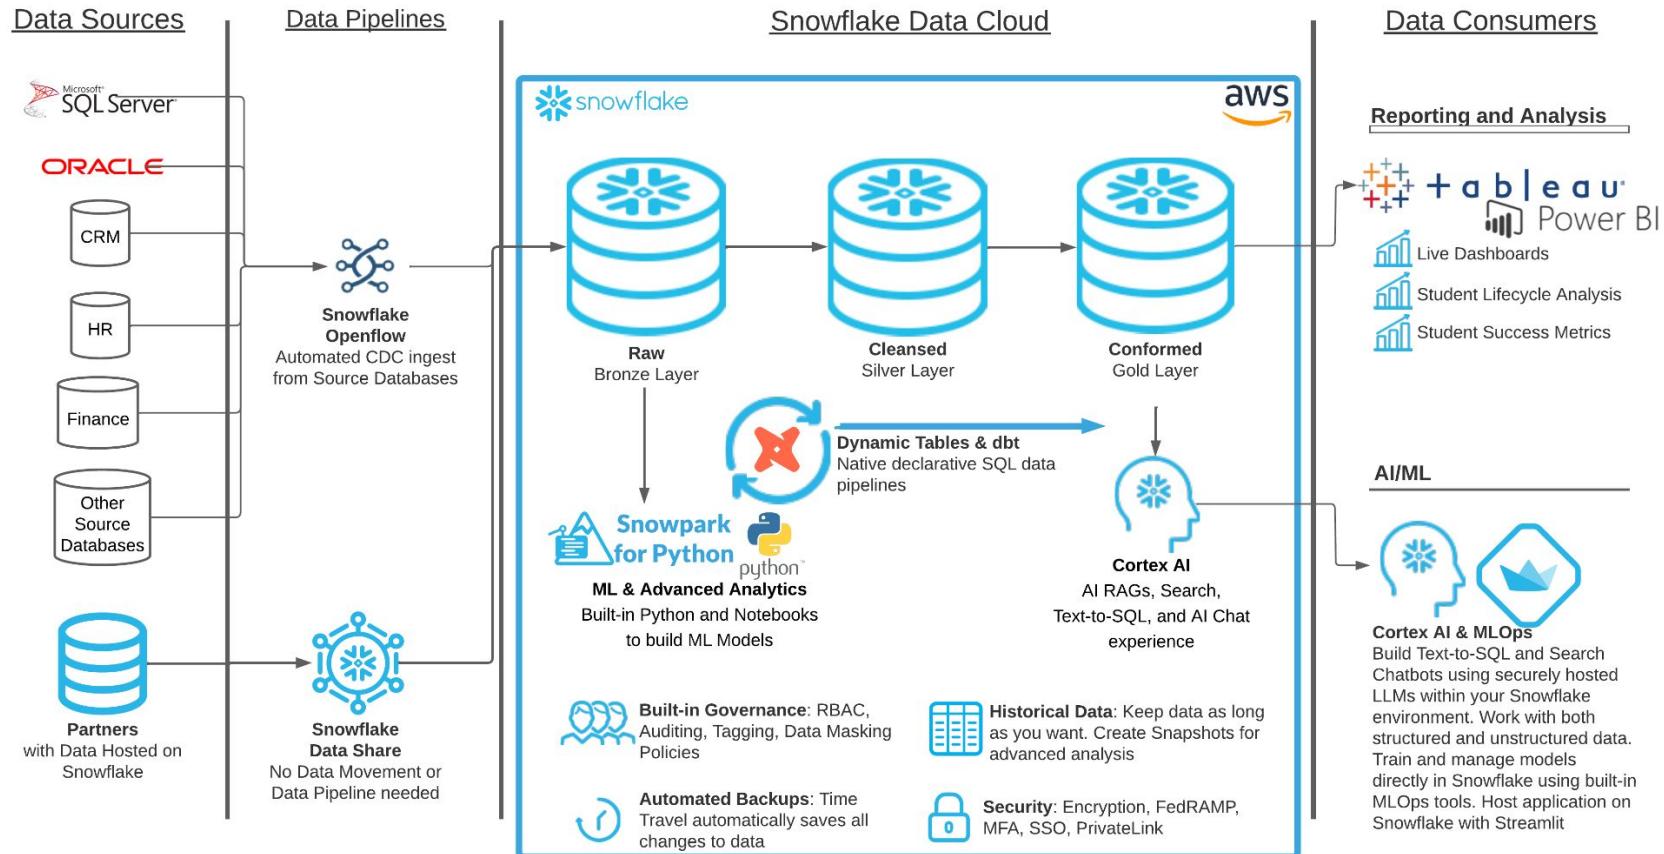


© 2023 Snowflake Inc. All Rights Reserved

UNIFIED GOVERNANCE



Example Data Architecture



Hands-on Lab: Data Exploration, Querying, & Analytics

1:30 PM - 3:30 PM



© 2025 Snowflake Inc. All Rights Reserved

The Fundamentals: Getting Started with Data



What is a Relational Database?

- Collection of data that organizes data in tables and maintains relationships between the tables for information retrieval
- Enterprise RDBMS include:
 - Microsoft SQL Server
 - Oracle Database
 - Oracle MySQL
 - IBM DB2
 - SAP Sybase
 - Teradata
 - PostgreSQL
 - **SNOWFLAKE!!!**



Key Terms

- A **table** is the basic unit of data storage in a database. Data is stored in rows and columns. You can specify rules for each column of a table like establishing its data type, forcing a column to contain a value in every row, or ensuring unique records
- A **view** is a tailored presentation of data selected from one or more tables, possibly including other views. A view contains no actual data but rather derives what it shows from the tables and views on which it is based. Therefore, a view can be thought of as a virtual table
- A **schema** is like a folder within a database similar to a directory on a file system that groups tables together
- The **join** between tables is made by a **Primary Key – Foreign Key** pair, where the Foreign Key in a table is the Primary Key of another table
- A **query** is used to create, manipulate, or view database objects
- **SQL** stands for Structured Query Language as is sometimes pronounced “sequel”



SELECT

- A SELECT statement is used to fetch records from a database table in the form of a table called a result set
- One or more columns may be specified, or the asterisk (*) can be used to return all records from the table
- Use the LIMIT keyword to return only the first n rows from the result set
- The result set is not stored in the database without the use of additional functions

SELECT column1, column2, column3 **FROM** table_name

SELECT * FROM table_name

SELECT column1, column2, column3 **FROM** table_name **LIMIT** 5



LIMIT

- The **LIMIT** keyword can be added to a **SELECT** statement to limit the number of records that are returned
- This is a best practice when developing and testing out queries to have quicker return times and to avoid higher compute costs
- Remember to remove the **LIMIT** and try the queries completely before productionizing them

```
SELECT column1, column2, column3 FROM table_name LIMIT 100
```



WHERE

- The **WHERE** clause can be added to a **SELECT** statement to filter records that meet certain conditions, i.e. return a subset of the table
- We use expressions to define these conditions
 - Most conditional operators are self-explanatory (=, <, >, etc.)
 - A common operator for “not equals” is \neq or !=

```
SELECT column1, column2, column3 FROM table_name WHERE column1 = 'ABC'
```

```
SELECT column1, column2, column3 FROM table_name WHERE column2 > 100
```

```
SELECT column1, column2, column3 FROM table_name WHERE column3  $\neq$  0
```



AND/OR

- An **AND/OR** operator can be added to a **WHERE** clause to filter records based on multiple conditions
- The **AND** operator returns all records where both conditions are true
- The **OR** operator returns all records if either the first condition or the second condition are true

```
SELECT column1, column2, column3 FROM table_name
```

```
WHERE a_condition AND another_condition;
```

```
SELECT column1, column2, column3 FROM table_name
```

```
WHERE a_condition OR another_condition;
```



IN

- The **IN** operator can be used to specify multiple values in a **WHERE** clause
- This is equivalent to a series of **OR** clauses

SELECT column1, column2, column3 **FROM** table_name

WHERE column1 **IN** (value1, value2, ..., valueN)

SELECT column1, column2, column3 **FROM** table_name

WHERE column1 = value1 **OR** column1 = value **OR** ... **OR** column1 = valueN



BETWEEN

- The **BETWEEN** operator can be used in a **WHERE** clause to select values within a range

```
SELECT column1, column2, column3 FROM table_name
```

```
WHERE column2 BETWEEN value1 AND value2
```



LIKE

- The **LIKE** operator can be used in a **WHERE** clause to search for a specified pattern in a column
- **LIKE** supports the % wildcards

```
SELECT column1, column2, column3 FROM table_name
```

```
WHERE column2 LIKE '%value1%'
```



NULL / NOT NULL

- **NULL** means that a value has not been entered for a particular column in a row, i.e. the value is missing
- **NULL** is not numeric or any other value so it can't equal itself or anything else
- **NOT** creates the negation of an expression
- To test for a **NULL** value you will want to use the expressions **IS NULL** or **IS NOT NULL**

```
SELECT column1, column2 FROM table_name
```

```
WHERE column2 IS NULL
```

```
SELECT column1, column2 FROM table_name
```

```
WHERE column2 IS NOT NULL
```



ORDER BY

- The **ORDER BY** keyword can be added to a **SELECT** statement to sort the records
- The typical default order is ascending (**ASC**), but you can also specify descending (**DESC**)

```
SELECT column1, column2, column3 FROM table_name
```

```
ORDER BY column2, column1, column3 DESC
```



Aggregation and Summarization: Finding the Insights



Aggregate Functions

- The following functions, including **COUNT**, calculate totals or statistics by column
- These statistics can apply to all rows or groupings of rows using **GROUP BY**
- Other SQL implementations may include additional statistics (standard deviation, variance, median, etc.)

Function	Description
COUNT()	Counts the total number of records
SUM()	Sum of numeric values
AVG()	Finds the average
MIN()	Returns the smallest number
MAX()	Returns the largest number



GROUP BY

The SQL **GROUP BY** clause is used in collaboration with the SELECT statement to aggregate by groups (similar to Excel's pivot table).

The **GROUP BY** clause follows the **WHERE** clause and precedes the **ORDER BY** clause (if present)

This clause regularly makes use of SQL's aggregate functions; when used with **GROUP BY**, each aggregate function produces a single value for each group

```
SELECT column1, column2, count(1) as ct, sum(column1) as sum1
```

```
FROM table_name
```

```
WHERE a_condition
```

```
GROUP BY column1, column2
```

```
ORDER BY column2, column1
```



HAVING

- The **HAVING** clause can be added to **GROUP BY** to filter records that meet certain conditions, i.e. return a subset of the groupings
- Operators are the same as the **WHERE** clause
-

SELECT column1, column2, count(1) as counts **FROM** table_name

GROUP BY column1, column2 **HAVING** counts > 5

SELECT column1, count(1) as counts, sum(column2) as sum2

FROM table_name **GROUP BY** column1

HAVING counts > 1 and sum2 > 10000



DISTINCT

- The **DISTINCT** function returns the unique values in a column or unique records across multiple columns
- The combination of **COUNT** and **DISTINCT** returns the number of unique values

SELECT DISTINCT column1 **FROM** table_name

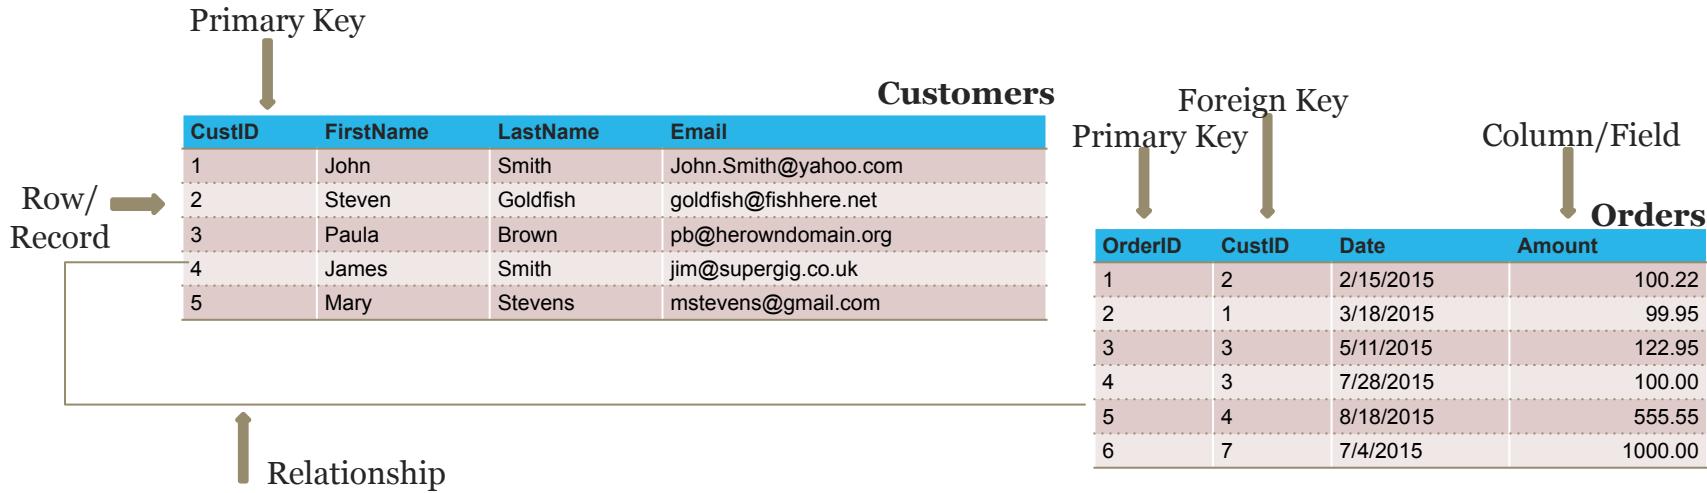
SELECT COUNT(DISTINCT column1) **FROM** table_name



The Relational Power: Joining Data Together



Relational Tables



Primary Key

- Primary key is a unique identifier of records in a table
- Primary key values may be generated manually or automatically
- A primary key can consist of more than one field

Athlete Table				
ID	Name	Country	Phone Number	E-mail Address
3245	Mo Farah	UK	02890 987001	mobot@ukathletics.co.uk
4532	Usain Bolt	Jamaica	02890 945002	usain.bolt@jamaica.com
7832	Jessica Ennis	UK	02890 922003	jennis@ukathletics.co.uk
2156	Michael Phelps	USA	02890 931004	michael.phelps@usathletics.com
9875	Yang Sun	China	02890 969005	yang.sun@swimming.cn



Foreign Key

Primary key field

Parent table

ATHLETES				
Athlete ID	Name	Date of Birth	Place of Birth	
3245	Mo Farah	23/03/1983	Mogadishu, Somalia	
7832	Jessica Ennis	30/11/1987	Sheffield, England	
2156	Michael Phelps	30/06/1985	Baltimore, USA	

EVENTS				
Event ID	Event	Athlete ID	Sport	
95432	5,000m	3245	Athletics	
14564	10,000m	3245	Athletics	
65941	Heptathlon	7832	Athletics	
10341	100m Medley	2156	Swimming	
10354	100m Backstroke	2156	Swimming	

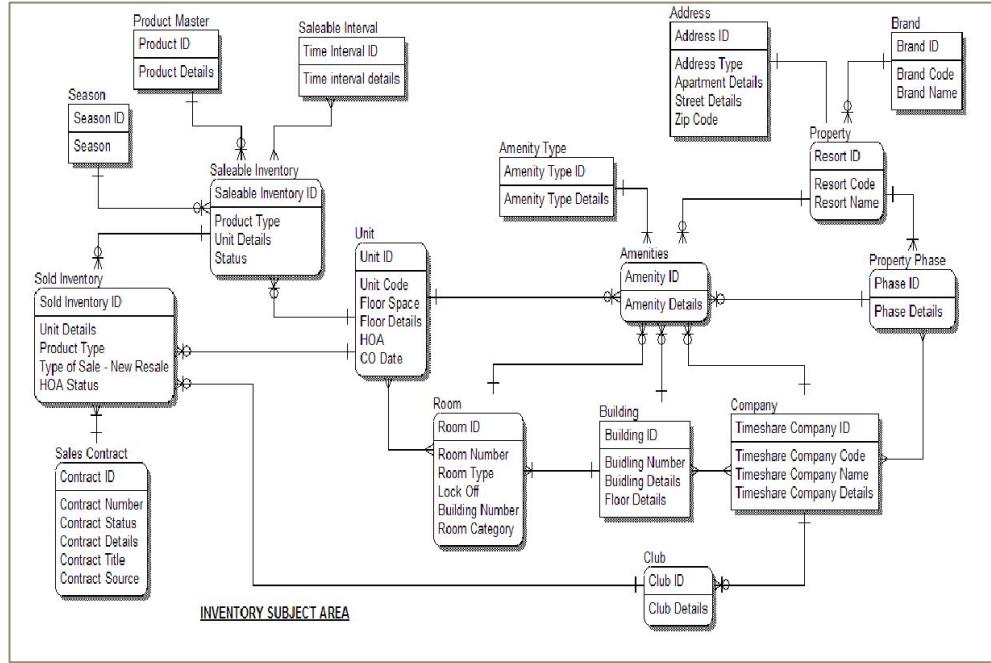
Relationship

Foreign key

Child table



Entity Relationship Diagram



JOINS

A **JOIN** combines fields from two tables using a key column shared by the tables

```
SELECT table_a.column1, table_a.column2, table_b.column3
```

```
FROM table_a JOIN table_b ON table_a.column1 = table_b.column1
```



Join Types

	INNER JOIN Selects rows from both tables as long as there is a match between the key columns in both tables
	LEFT JOIN Returns all rows from the left table (1), with the matching rows in the right table (2); the result is NULL in the right side when there is no match
	RIGHT JOIN Returns all rows from the right table (2), with the matching rows in the left table (1); the result is NULL in the left side when there is no match
	FULL OUTER JOIN Returns all rows from the left table (1) and from the right table (2); the result is NULL in the left or right side when there is no match

CASE STATEMENTS

- **CASE** statements allow conditional logic in SQL queries.
- They evaluate conditions and return a result when the first condition is met

```
SELECT  
    CASE
```

```
        WHEN column1 = 1 THEN 'one'  
        WHEN column1 = 2 THEN 'two'  
        WHEN column1 = 3 THEN 'three'  
        ELSE 'zero'
```

```
    END AS number_col
```

```
FROM table_name
```



Data Management & Reporting



Uploading Data through Snowflake UI

Load Data into Table • COMPUTE_WH

Drag and drop to upload files
or
[Browse](#)
or
[Add from Stage](#)

File size limit: 250MB
Supported formats: CSV/TSV, json, orc, avro, parquet

Select or create a database and schema

Schema **BRONZE.ADMIN** ▼ [+ Database](#)

Select or create a table

[Cancel](#) [Back](#) [Next](#)



Views

- A **VIEW** allows the result of a query to be accessed as if it were a table.
- The **SELECT** query is specified in the **CREATE VIEW** statement.

CREATE OR REPLACE VIEW view_name **AS**

SELECT

CASE

WHEN column1 = 1 **THEN** 'one'

WHEN column1 = 2 **THEN** 'two'

WHEN column1 = 3 **THEN** 'three'

ELSE 'zero'

END AS number_col,
column2, column3

FROM table_name



Snowflake Dynamic Tables

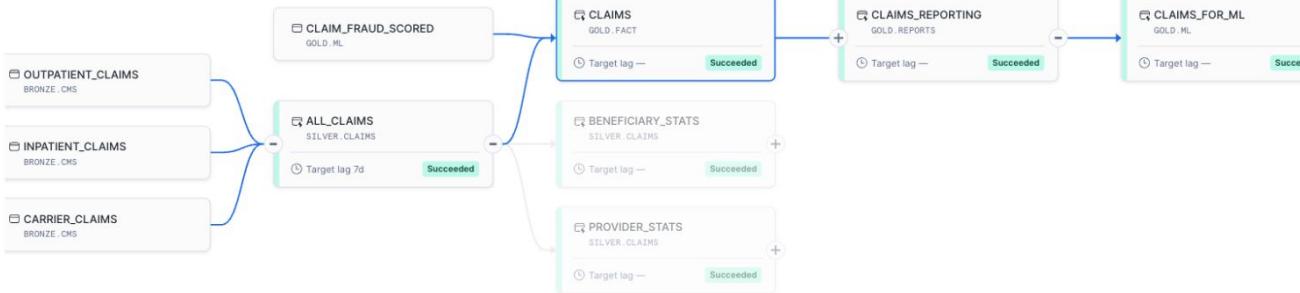


GOLD / FACT / CLAIMS

Dynamic Table SYSADMIN 6 days ago 112.5M COMPUTE_WH Incremental

Table Details Columns Data Preview Graph Refresh History

+ - () ☰ 🔍 | • COMPUTE_WH ⏪



CLAIMS

GOLD.FACT

Details Definition

Lag Metrics

Last Refresh

Succeeded

Time Wit... ⓘ

—

Current Lag

5d 19h 49m 1s

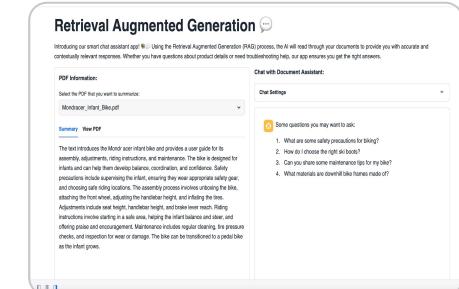
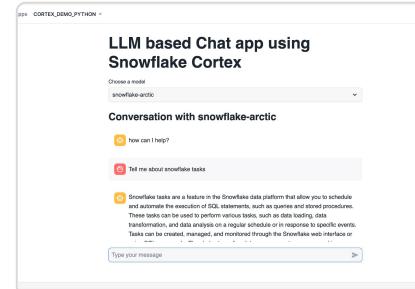
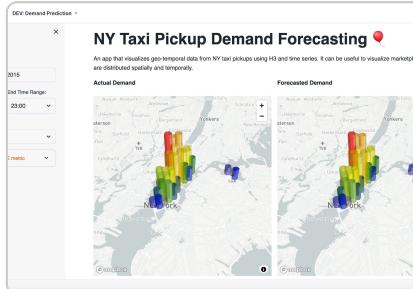
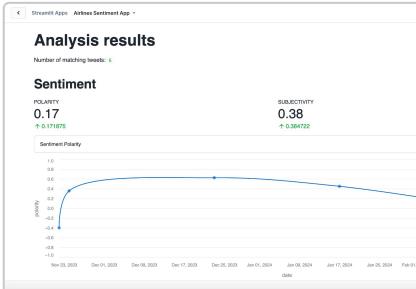
Target Lag

(downstream)

BONUS: Using AI to Build Dashboards



Streamlit-in-Snowflake



Natural Language Processing

Sentiment Analysis

NLP models that showcase text summarization, sentiment analysis, translation

Mapping & Geospatial

Interactive Maps

Build powerful geospatial experiences using H3 and Mapbox



Streamlit



© 2025 Snowflake Inc. All Rights Reserved

LLM Integration

Chat-based Applications

Applications which leverage leading LLM models from Snowflake Cortex

RAG Applications

Retrieval Augmented Gen

Ask questions of data siloed in documents using RAG with Snowflake Cortex